



Lab 7.1 - Timeouts and Retries with Linkerd

Overview

In this lab, you'll be configuring timeouts and retries for service-to-service networking with Linkerd. Timeouts and retries are important for mitigating cascading failures when a microservice is not responding as expected.

1. Connect to the VM that hosts your Kubernetes clusters.
2. To start the cluster that contains the Linkerd service mesh, issue this command:

```
yourname@ubuntu-vm:~$ docker start $(docker ps -a -f  
name=linkerd-control-plane -q)
```

```
8af4c08524df
```

3. To switch the `kind` context to the Linkerd cluster, use this command:

```
yourname@ubuntu-vm:~$ kubectl config use-context kind-linkerd
```

```
Switched to context "kind-linkerd".
```

4. In the previous labs you used the Emojivoto application to demonstrate basic service mesh functionality. Since Linkerd is currently unable to configure retries on gRPC requests, which is what Emojivoto uses, you will use a different example application for this lab. Run the following command to remove the Emojivoto application from your cluster. Note that it may take this command a few minutes to complete executing and return you to a command line prompt.

```
yourname@ubuntu-vm:~$ kubectl delete namespace emojiivoto
```

```
namespace "emojiivoto" deleted
serviceaccount "emoji" deleted
serviceaccount "voting" deleted
serviceaccount "web" deleted
service "emoji-svc" deleted
service "voting-svc" deleted
service "web-svc" deleted
deployment.apps "emoji" deleted
deployment.apps "vote-bot" deleted
deployment.apps "voting" deleted
deployment.apps "web" deleted
```

5. Make sure Linkerd and Nginx Ingress have started and have all resources ready before continuing. The following command will list the `Pods` in all namespaces. All should have a status of `Running` and a ready status with the same numbers on the left and right sides of the slash ("/"). If any of the lines don't show those statuses, wait a few minutes and reissue the command again.

```
yourname@ubuntu-vm:~$ kubectl get pods -n linkerd
```

NAMESPACE	NAME			
linkerd-destination-7df897cf9b-4w49m		4/4	Running	8
(2m42s ago)	20h			
linkerd-identity-5f4dbf785d-vxnrm		2/2	Running	5
(2m42s ago)	20h			
linkerd-proxy-injector-5497f6fbd7-qzwvx		2/2	Running	4
(2m42s ago)	20h			

```
yourname@ubuntu-vm:~$ kubectl get pods -n ingress-nginx
```

ingress-nginx-admission-create-r7gvr	0/1	Completed	0
40h			
ingress-nginx-admission-patch-6prps	0/1	Completed	0
40h			
ingress-nginx-controller-b66cc4b74-pp4kh	1/1	Running	4
(2m7s ago) 40h			

6. For this lab, you will use another application provided by Buoyant (the creators of Linkerd) to demonstrate Linkerd's ability to perform retries and timeouts on individual services in the mesh. Install the [booksapp](#) application:

```
yourname@ubuntu-vm:~$ kubectl create ns booksapp && \
  curl -sL https://run.linkerd.io/booksapp.yml \
  | linkerd inject - \
  | kubectl -n booksapp apply -f -
```

```
namespace/booksapp created
```

```
service "webapp" skipped
deployment "webapp" injected
service "authors" skipped
deployment "authors" injected
service "books" skipped
deployment "books" injected
deployment "traffic" injected
```

```
service/webapp created
deployment.apps/webapp created
service/authors created
deployment.apps/authors created
service/books created
deployment.apps/books created
deployment.apps/traffic created
```

7. Next, configure the Nginx Ingress to use the new app.

```
yourname@ubuntu-vm:~$ kubectl apply -f - <<EOF
---
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  namespace: booksapp
  name: ingress-booksapp
  annotations:
    ingress.kubernetes.io/rewrite-target: /
    nginx.ingress.kubernetes.io/service-upstream: "true"
    consul.hashicorp.com/connect-inject: "true"
spec:
  ingressClassName: nginx
  rules:
```

```
- http:
  paths:
  - pathType: Prefix
    path: "/"
    backend:
      service:
        name: webapp
        port:
          number: 7000
EOF
```

```
ingress.networking.k8s.io/ingress-booksapp created
```

8. Create a Kubernetes `port-forward`:

```
yourname@ubuntu-vm:~$ kubectl port-forward
service/ingress-nginx-controller -n ingress-nginx 8080:80
```

9. Navigate to 127.0.0.1:8080 in a web browser to interact with the new application. This application is a database of books and authors. Clicking a book will give you information about the book like the page count and author. Clicking an author will give you a list of the books they have written. You can also add authors and books to the database. Add several books to the database, and you will find that the function fails around half the time.

Add a Book

Title	<input type="text" value="Book1"/>
Author	<input type="text" value="James, P.D."/> ▼
Page count	<input type="text" value="10"/>
<input type="button" value="Add Book"/>	

Internal Server Error

10. Linkerd exposes the [ServiceProfile](#) configuration to see per-route metrics, and to configure timeouts and automatic retries on those routes. This allows for you to pinpoint

some errors and mitigate them at the network level while you work on patching the code. For HTTP services, there are two mechanisms for creating **ServiceProfiles**. You can use live traffic to [automatically create ServiceProfiles](#) based on requests through the proxy or you can use [OpenAPI \(Swagger\)](#) documentation for the service to generate a **ServiceProfile**. To guarantee you have a full understanding of the available routes, you will use the available OpenAPI docs for this service.

```
yourname@ubuntu-vm:~$ curl -sL
https://run.linkerd.io/booksapp/webapp.swagger \
  | linkerd -n booksapp profile --open-api - webapp \
  | kubectl -n booksapp apply -f -

serviceprofile.linkerd.io/webapp.booksapp.svc.cluster.local
created
```

11. Wait a few minutes, then take a look at the metrics for these routes from the command line. The **linkerd routes** command returns a snapshot of the performance of each route in the **ServiceProfile** above. The output appears to indicate that **POST** requests to the **/books** endpoint fail around half the time.

```
yourname@ubuntu-vm:~$ linkerd viz routes -n booksapp
service/webapp
```

ROUTE		SERVICE	SUCCESS	RPS	
LATENCY_P50	LATENCY_P95	LATENCY_P99			
GET /		webapp	100.00%	0.6rps	15ms
20ms	20ms				
GET /authors/{id}		webapp	100.00%	0.6rps	8ms
10ms	10ms				
GET /books/{id}		webapp	100.00%	1.1rps	8ms
18ms	20ms				
POST /authors		webapp	100.00%	0.6rps	9ms
19ms	20ms				
POST /authors/{id}/delete		webapp	100.00%	0.6rps	15ms
20ms	20ms				
POST /authors/{id}/edit		webapp	-	-	-
-	-				
POST /books		webapp	45.83%	3.0rps	
9ms	18ms				
20ms					
POST /books/{id}/delete		webapp	100.00%	0.6rps	8ms
10ms	10ms				

POST /books/{id}/edit	webapp	42.35%	1.4rps	16ms
84ms 97ms				
[DEFAULT]	webapp	-	-	-
-				

12. You now know that requests from the webapp service to the books service appear to be failing. You can get more visibility to this by creating a **ServiceProfile** for the upstream services as well. You can do this with OpenAPI documentation for the books service.

```
yourname@ubuntu-vm:~$ curl -sL
https://run.linkerd.io/booksapp/books.swagger \
| linkerd -n booksapp profile --open-api - books \
| kubectl -n booksapp apply -f -

serviceprofile.linkerd.io/books.booksapp.svc.cluster.local
created
```

13. Wait a few minutes, then run the **linkerd routes** command below to review the requests from the webapp service to the books service. You can see that **POST** and **PUT** actions seem to be the ones causing errors.

```
yourname@ubuntu-vm:~$ linkerd viz routes -n booksapp
deploy/webapp --to service/books
```

ROUTE	SERVICE	SUCCESS	RPS
LATENCY_P50 LATENCY_P95 LATENCY_P99			
DELETE /books/{id}.json	books	100.00%	1.3rps 5ms
9ms 10ms			
GET /books.json	books	100.00%	1.9rps 2ms
3ms 4ms			
GET /books/{id}.json	books	100.00%	2.5rps 2ms
3ms 17ms			
POST /books.json	books	44.25%	2.9rps 5ms
16ms 19ms			
PUT /books/{id}.json	books	54.05%	1.2rps 6ms
84ms 97ms			
[DEFAULT]	books	-	-
-			

14. It is a little strange that two independent routes, both the **POST** and **PUT**, are erroring at the same rate. You can dig a little deeper by seeing if both of these routes rely on

another service. Checking the output of `linkerd tap` with the command below will give you a log of the requests flowing through the books service. Hit `Control-C` to stop the log from scrolling so you can read it. From that log, you should find that the books service is making a lot of `HEAD` requests to the authors service. Based on the number of 503s that are being returned, those requests appear to be failing quite frequently.

```
yourname@ubuntu-vm:~$ linkerd viz tap -n booksapp deploy/books
```

```
req id=9:49 proxy=out src=10.244.0.53:37820 dst=10.244.0.50:7001
tls=true :method=HEAD :authority=authors:7001
:path=/authors/3252.json
rsp id=9:49 proxy=out src=10.244.0.53:37820 dst=10.244.0.50:7001
tls=true :status=503 latency=2197µs
end id=9:49 proxy=out src=10.244.0.53:37820 dst=10.244.0.50:7001
tls=true duration=16µs response-length=0B
```

15. Let's investigate this further. To start, create a `ServiceProfile` for the authors service from the swagger doc.

```
yourname@ubuntu-vm:~$ curl -sL
https://run.linkerd.io/booksapp/authors.swagger \
  | linkerd -n booksapp profile --open-api - authors \
  | kubectl -n booksapp apply -f -

serviceprofile.linkerd.io/authors.booksapp.svc.cluster.local
created
```

16. After creating the `ServiceProfile`, wait a few minutes and then check `linkerd routes` for requests from the books service to the authors service. You will see that the `HEAD` requests are failing.

```
yourname@ubuntu-vm:~$ linkerd viz routes -n booksapp deploy/books
--to service/authors
```

ROUTE	SERVICE	SUCCESS	RPS
LATENCY_P50 LATENCY_P95 LATENCY_P99			
DELETE /authors/{id}.json	authors	-	-
-			
GET /authors.json	authors	-	-
-			
GET /authors/{id}.json	authors	-	-
-			

HEAD	/authors/{id}.json	authors	54.87%	3.8rps	
1ms	2ms	3ms			
POST	/authors.json	authors	-	-	-
-	-				
[DEFAULT]		authors	-	-	-
-	-				

17. From the output of the commands in the previous steps, you can see that linkerd is reporting similar behavior to what we have been seeing while interacting with the application. The actions that add books to the database fail around half the time. Based on your debugging work, it seems that this is actually due to an issue with the **HEAD** requests to the authors service. This issue needs to be remedied quickly. Since coding can be a time consuming activity, Linkerd gives you the ability to easily put some mitigations in place while you work on debugging what went wrong.

- a. The **ServiceProfile** is not just capable of *reporting* per-route metrics; it can also *configure* per-route rules for linkerd. Based on the data above, adding books to the database appears to be failing around half the time. You can configure automatic retries on those actions so that Linkerd will try to re-send the request instead of reporting an error to your users. Manually edit the **ServiceProfile** by running `kubectl edit` to open it in a text editor.

```
yourname@ubuntu-vm:~$ kubectl edit -n booksapp
serviceprofile authors.booksapp.svc.cluster.local
```

[Opens a text editor on success]

- b. Examine the **ServiceProfile** and find the **HEAD /authors/{id}.json** routes. Adding `isRetryable: true` to that configuration will tell Linkerd that it should retry requests to that route. Edit the **ServiceProfile** so it matches the `isRetryable: true` configuration below. Save and exit the editor to apply the edit.

```
...
- condition:
  method: HEAD
  pathRegex: /authors/[^/]*\.json
  isRetryable: true
  name: HEAD /authors/{id}.json
```

```
serviceprofile.linkerd.io/authors.booksapp.svc.cluster.local
edited
```


18. You have now configured retries for **HEAD** requests to `/authors/{id}.json` in the **authors** service. After giving it some time, check **linkerd routes** to confirm that the success rate has now climbed back up to 100%.

```
yourname@ubuntu-vm:~$ linkerd viz routes -n booksapp
service/webapp
```

ROUTE		SERVICE	SUCCESS	RPS	
LATENCY_P50	LATENCY_P95	LATENCY_P99			
GET /		webapp	100.00%	0.9rps	16ms
25ms	29ms				
GET /authors/{id}		webapp	100.00%	0.9rps	8ms
15ms	19ms				
GET /books/{id}		webapp	100.00%	1.7rps	8ms
10ms	10ms				
POST /authors		webapp	100.00%	0.9rps	13ms
19ms	20ms				
POST /authors/{id}/delete		webapp	100.00%	0.9rps	18ms
29ms	30ms				
POST /authors/{id}/edit		webapp	-	-	-
-	-				
POST /books		webapp	100.00%	1.7rps	15ms
20ms	28ms				
POST /books/{id}/delete		webapp	100.00%	0.9rps	8ms
10ms	10ms				
POST /books/{id}/edit		webapp	100.00%	0.9rps	18ms
72ms	94ms				
[DEFAULT]		webapp	-	-	-
-	-				

19. Congratulations! You have successfully mitigated an issue for your users without having to make any code changes. Your team now has time to debug and find the right changes to make, and does not have to rush to push out a bug fix!
20. While this is exciting for your users, digging in a little deeper will show you the effect it is having on your system. If you pass the `-o wide` flag to **linkerd routes**, it will show you that the **EFFECTIVE_SUCCESS** and **ACTUAL_SUCCESS** rates are quite different. In the sample output below, **EFFECTIVE_SUCCESS** is 100.00% but **ACTUAL_SUCCESS** is 48.14%. This is because the Linkerd sidecar proxy in the books service is now sending more requests to the authors service when a request fails. You can see this in the

EFFECTIVE_RPS and **ACTUAL_RPS** values (2.6rps and 5.4 rps, respectively, in the sample output).

```
yourname@ubuntu-vm:~$ linkerd viz routes -n booksapp deploy/books
--to service/authors -o wide
```

ROUTE		SERVICE	EFFECTIVE_SUCCESS		
EFFECTIVE_RPS	ACTUAL_SUCCESS	ACTUAL_RPS	LATENCY_P50		
LATENCY_P95	LATENCY_P99				
DELETE /authors/{id}.json		authors		-	
-	-	-	-	-	-
-					
GET /authors.json		authors		-	
-	-	-	-	-	-
-					
GET /authors/{id}.json		authors		-	
-	-	-	-	-	-
-					
HEAD /authors/{id}.json		authors	100.00%		
2.6rps	48.14%	5.4rps	3ms	10ms	
17ms					
POST /authors.json		authors		-	
-	-	-	-	-	-
-					
[DEFAULT]		authors		-	
-	-	-	-	-	-
-					

21. The other effect this has had is it slightly increases the latency of adding books. Running `linkerd routes` will show that there is a relatively high P95 latency.

```
yourname@ubuntu-vm:~$ linkerd viz routes -n booksapp
deploy/webapp --to svc/books
```

ROUTE		SERVICE	SUCCESS	RPS	
LATENCY_P50	LATENCY_P95	LATENCY_P99			
DELETE /books/{id}.json		books	100.00%	0.9rps	8ms
10ms	10ms				
GET /books.json		books	100.00%	1.8rps	3ms
5ms	5ms				
GET /books/{id}.json		books	100.00%	2.6rps	2ms
3ms	3ms				

POST	/books.json	books	100.00%	1.6rps	14ms
	19ms	20ms			
PUT	/books/{id}.json	books	100.00%	0.8rps	16ms
	82ms	97ms			
[DEFAULT]		books	-	-	-
-	-				

22. This issue, like retries, can be mitigated by Linkerd until your team can fully resolve the issue.

- a. Run `kubectl edit` to open the `ServiceProfile` in a text editor.

```
yourname@ubuntu-vm:~$ kubectl edit -n booksapp
serviceprofile books.booksapp.svc.cluster.local
```

```
[Opens a text editor on success]
```

- b. Find the `PUT` route you need to configure retries on, and add the `timeout: 25ms` configuration below. This will cause attempts to time out more quickly. Save and exit the editor to apply the edit.

```
...
- condition:
  method: PUT
  pathRegex: /books/[^/]*\.json
  name: PUT /books/{id}.json
  timeout: 25ms

serviceprofile.linkerd.io/books.booksapp.svc.cluster.local
edited
```

23. Wait a few minutes and run `linkerd routes` again to see the latest data. The latency numbers include time spent in the webapp itself, so they will remain above the 25 ms threshold you set. You can see that the timeouts are working because the **EFFECTIVE_SUCCESS** rate drops as the linkerd proxy times out on some requests. In the example output below, that rate is 92.16%. There are always tradeoffs when mitigating issues with microservices, and you must decide if setting around a 90% success rate with moderate latency is a reasonable way to leave the system while your team works on a fix for the service.

```
yourname@ubuntu-vm:~$ linkerd -n booksapp viz routes
deploy/webapp --to svc/books -o wide
```

ROUTE	SERVICE		EFFECTIVE_SUCCESS	
EFFECTIVE_RPS	ACTUAL_SUCCESS	ACTUAL_RPS	LATENCY_P50	
LATENCY_P95	LATENCY_P99			
DELETE /books/{id}.json	books	100.00%		
0.9rps	100.00%	0.9rps	8ms	10ms
10ms				
GET /books.json	books	100.00%		
1.9rps	100.00%	1.9rps	3ms	4ms
5ms				
GET /books/{id}.json	books	100.00%		
2.7rps	100.00%	2.7rps	2ms	3ms
3ms				
POST /books.json	books	100.00%		
1.6rps	100.00%	1.6rps	14ms	19ms
20ms				
PUT /books/{id}.json	books	92.16%		
0.8rps	100.00%	0.8rps	11ms	77ms
96ms				
[DEFAULT]	books	-		
-	-	-	-	-

24. You've reached the end of this lab. You will continue to use the Linkerd cluster in your next lab, so leave it running if you plan to move on to the next lab now. Otherwise, stop the Linkerd cluster by running this command:

```
yourname@ubuntu-vm:~$ docker stop $(docker ps -a -f
name=linkerd-control-plane -q)
```

```
8af4c08524df
```